

Assignment-4

SMS Spam Classification

Assignment Date	30 September 2022
Student Name	Guru Hari Venkat S
Student Roll Number	513419106012
Maximum Marks	2 Marks

SMS SPAM Classification

Import required libraries

```
import pandas as pd
import numpy as np
import re
import collections
import contractions
import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('dark_background')
import nltk
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
import warnings
warnings.simplefilter(action='ignore', category=Warning)
import keras
from keras.layers import Dense, Embedding, LSTM, Dropout
from keras.models import Sequential
from keras.preprocessing.text import Tokenizer
from keras_preprocessing.sequence import pad_sequences
```

Download the dataset

```
df = pd.read_csv("spam.csv", encoding='ISO-8859-1')
df.shape
```

```
(5572, 5)
```

Read dataset and do pre-processing

```
df.head(10)
```

```
      v1  ... Unnamed: 4
0   ham  ...         NaN
1   ham  ...         NaN
2  spam  ...         NaN
```

```

3    ham    ...    NaN
4    ham    ...    NaN
5   spam    ...    NaN
6    ham    ...    NaN
7    ham    ...    NaN
8   spam    ...    NaN
9   spam    ...    NaN

```

```
[10 rows x 5 columns]
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0    v1              5572 non-null   object
1    v2              5572 non-null   object
2    Unnamed: 2      50 non-null     object
3    Unnamed: 3      12 non-null     object
4    Unnamed: 4      6 non-null      object
dtypes: object(5)
memory usage: 217.8+ KB

```

```
df.isnull().any()
```

```

v1          False
v2          False
Unnamed: 2   True
Unnamed: 3   True
Unnamed: 4   True
dtype: bool

```

```
df.isnull().sum()
```

```

v1          0
v2          0
Unnamed: 2   5522
Unnamed: 3   5560
Unnamed: 4   5566
dtype: int64

```

```

df.drop(["Unnamed: 2", "Unnamed: 3", "Unnamed: 4"], axis=1,
inplace=True)
df.columns = ["Spam/Ham", "Message"]

```

```
df.head()
```

```

      Spam/Ham      Message
0      ham  Go until jurong point, crazy.. Available only ...
1      ham                Ok lar... Joking wif u oni...

```

```
def preprocessing(data):
    sms = contractions.fix(data)
    sms = sms.lower()
    sms = re.sub(r'https?://S+|www.S+', "", sms).strip()
    sms = re.sub("[^a-z ]", "", sms)
    return sms

X = df["Message"].apply(preprocessing)

from sklearn.preprocessing import LabelEncoder
lb_enc = LabelEncoder()
y = lb_enc.fit_transform(df["Spam/Ham"])

tokenizer = Tokenizer()
tokenizer.fit_on_texts(X)
text_to_sequence = tokenizer.texts_to_sequences(X)

max_length_sequence = max([len(i) for i in text_to_sequence])
padded_sms_sequence = pad_sequences(text_to_sequence,
maxlen=max_length_sequence,
padding = "pre")
```

```
array([[ 0, 0, 0, ..., 50, 3915, 134],
       [ 0, 0, 0, ..., 419, 1, 1715],
       [ 0, 0, 0, ..., 2653, 348, 2654],
       ...,
       [ 0, 0, 0, ..., 8472, 222, 8473],
       [ 0, 0, 0, ..., 144, 17, 45],
       [ 0, 0, 0, ..., 3, 61, 233]], dtype=int32)
```

Create Model

Add Layers (LSTM, Dense-(Hidden Layers), Output)

```
TOT_SIZE = len(tokenizer.word_index)+1
```

```
def create_model():
    lstm_model = Sequential()
    lstm_model.add(Embedding(TOT_SIZE, 32,
input_length=max_length_sequence))
    lstm_model.add(LSTM(100))
    lstm_model.add(Dropout(0.4))
    lstm_model.add(Dense(20, activation="relu"))
    lstm_model.add(Dropout(0.3))
    lstm_model.add(Dense(1, activation = "sigmoid"))
    return lstm_model
lstm_model = create_model()
```

Compile the Model

```
lstm_model.compile(loss = "binary_crossentropy", optimizer = "adam",
metrics = ["accuracy"])
```

```
lstm_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 172, 32)	271200
lstm (LSTM)	(None, 100)	53200
dropout (Dropout)	(None, 100)	0
dense (Dense)	(None, 20)	2020
dropout_1 (Dropout)	(None, 20)	0
dense_1 (Dense)	(None, 1)	21
Total params: 326,441		
Trainable params: 326,441		

Non-trainable params: 0

Fit the Model

```
lstm_model.fit(padded_sms_sequence, y, epochs = 5,  
validation_split=0.2, batch_size=16)
```

Epoch 1/5

```
279/279 [=====] - 34s 112ms/step - loss:  
0.2117 - accuracy: 0.9361 - val_loss: 0.0662 - val_accuracy: 0.9821
```

Epoch 2/5

```
279/279 [=====] - 34s 122ms/step - loss:  
0.1441 - accuracy: 0.9746 - val_loss: 0.0819 - val_accuracy: 0.9830
```

Epoch 3/5

```
279/279 [=====] - 32s 114ms/step - loss:  
0.0272 - accuracy: 0.9953 - val_loss: 0.0638 - val_accuracy: 0.9848
```

Epoch 4/5

```
279/279 [=====] - 31s 111ms/step - loss:  
0.0135 - accuracy: 0.9980 - val_loss: 0.0622 - val_accuracy: 0.9865
```

Epoch 5/5

```
279/279 [=====] - 31s 111ms/step - loss:  
0.0111 - accuracy: 0.9978 - val_loss: 0.0684 - val_accuracy: 0.9839
```

<keras.callbacks.History at 0x7f4de753af10>

Save The Model

```
lstm_model.save('sms.h5')
```

Test The Model

```
def predict_spam(predict_msg):  
    new_seq = tokenizer.texts_to_sequences(predict_msg)  
    padded = pad_sequences(new_seq, maxlen =max_length_sequence,  
                           padding = 'pre')  
    return (lstm_model.predict(padded))
```

```
message = str(input('Enter your message: '))
```

```
pred = predict_spam(message)
```

```
print(f"Prediction: {pred[0][0]}\n")
```

```
print(f"The sms is a {'spam' if pred[0][0]<0.5 else 'ham'}")
```

Enter your message: you've won something, the IRS is trying to contact
you, you have a refund coming or asking you to verify your bank
account

```
4/4 [=====] - 0s 36ms/step
```

```
Prediction: 0.005271350499242544
```

The sms is a spam